

Gold Pipeline Run Report

Session: gold_20260711_162915_essay_001_874887d6 | Student: student_123 | Config: gold_engine_commands_full_v1_4_10.json (det_vip LLM detector, gpt-4o-mini)

Overall Result

qa_status: passed — all 28 pipeline stages completed successfully with zero QA gate issues, zero missing/invalid artifacts, zero boundary issues. This is the first real run of the new det_vip LLM-based detector, replacing the placeholder regex-only detector.

Detector (det_vip_v18d_2.py via det_vip_cli_bridge_v1.py)

Metric	Value
LLM status	enabled (gpt-4o-mini)
LLM calls attempted / succeeded	32 / 18
Estimated LLM cost	\$0.0069
Raw candidates	63
Chargeable (scoreable) rows	15
Review-only rows (correctly excluded)	22
spaCy status	unavailable
LanguageTool status	unavailable
Resource quality status	ready

Bug-fix verification: the review-only vs. chargeable fix (aliasing chargeable across every row list det_vip produces, not just student_rows) held correctly at real volume — all 22 review-only rows were verified as chargeable=False, none leaked through as scoreable evidence.

Two things worth a closer look: (1) spaCy and LanguageTool both reported "unavailable" inside the actual pipeline run, despite both working when you tested them standalone in PowerShell moments earlier — likely an environment difference between your interactive shell and the subprocess environment the orchestrator launches (e.g. LanguageTool's Java dependency or spaCy's model path not resolving the same way inside a spawned subprocess). Detection still worked via LLM + rule passes, but you're not yet getting the full rule+spaCy+LanguageTool+LLM signal det_vip was designed for. (2) 14 of 32 LLM calls failed (44%) with no error detail captured — worth monitoring if this recurs, since it represents lost detection signal and wasted spend.

Downstream Chain

Stage	Result
Errormap	15 errors, all resolved to a known criterion/capacity_domain (0 unknown)
→ by criterion	grammar: 10, lexical_resource: 5
→ by capacity_domain	sentence_control: 10, lexical_precision: 5

Score Contract	Band 5.0 overall (TR 5 / CC 5 / LR 5 / GRA 5), confidence: normal
Verifier / Adjudicator	pass / confirmed
Progress Tracker	Correct identity; score_events: 4 (accumulating correctly)
Continuity (LIE)	cold_start: false, sessions_analyzed: 4
LRET	canonical resources loaded; qa_status: ok; 3 fix, 2 enhance, 18 clarify, 31 keep units
Writing Coach	Selected move: "Article Meaning to Noun Phrase" (Article Control)
Practice Session	1 exercise built
Logs	No errors or unexpected warnings in any of the 28 stage logs

Recommended Next Steps

- Investigate why spaCy/LanguageTool show unavailable inside the pipeline subprocess despite working standalone -- check whether the orchestrator's Python interpreter/PATH/working directory matches your interactive shell's.
- Monitor the 44% LLM call failure rate on future runs; if it persists, det_vip's LLMTracker error details may need surfacing to diagnose (rate limits vs. timeouts vs. response-format issues).
- Practice session produced only 1 exercise this run -- worth a quick sanity check against the directive's focus areas if that seems low for a 10-minute session.